An evidence-controlled Python agent

The model chooses the next evidence action.
Python checks its scope and renders the supported answer.

39 curated facts

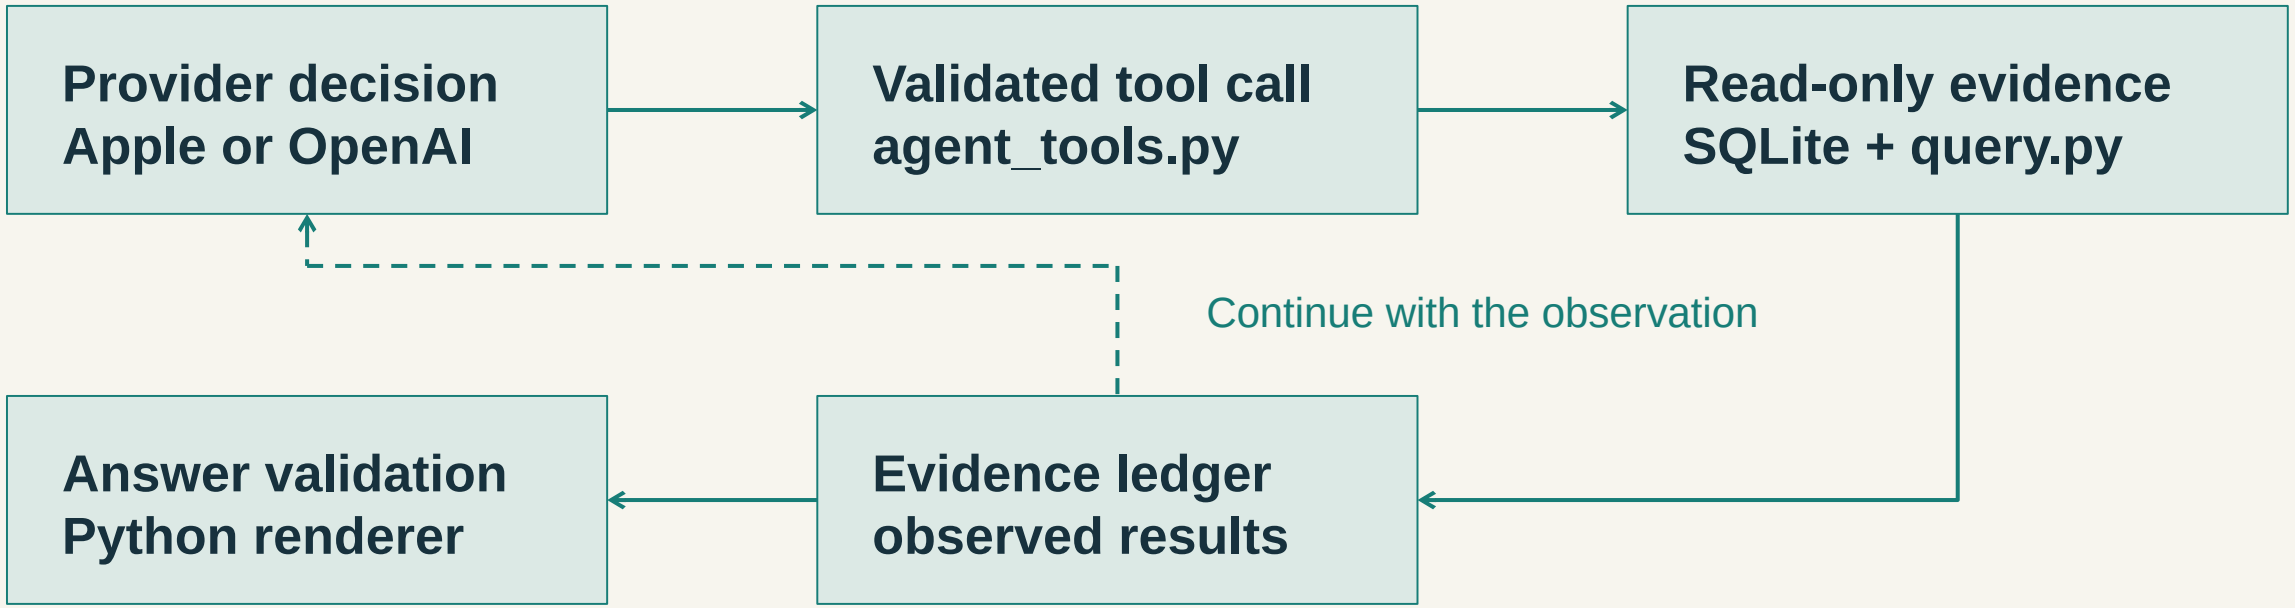
5 fixed tools

1 standalone repo

The agentic loop

One structured decision at a time, with checked observations returned to the model

search_facts, get_fact, compare_facts, get_source, coverage_gap



Finish only with checked references

Session state retains topic, geography, period and entity. /new clears it.
The host exposes tool errors, unresolved scope and bounded-loop failures.

Modules and repository structure

Repository path	Responsibility
agent.py / agent_loop.py	CLI, bounded loop, session state, evidence ledger and rendering
providers.py apple_provider.py / apple_worker.py	OpenAI Responses or optional Apple SDK adapter with an isolated worker
agent_tools.py / agent_protocol.py contracts/ / contract_validation.py	Five typed tools, strict decisions, comparison safeguards and contract validation
evidence/london_office_evidence_pack/	Curated JSON, evidence.sqlite and query.py for deterministic access and rebuild
scripts/ / tests/ / eval/	Integrity checks, offline integration, six-case replay and human review
docs/user/ / agent_skills/ / slides/	User guides, task guidance, editable deck and presentation artifacts

Evidence enforcement and provenance

Curated inputs and controlled access

sources.json / facts.json / coverage.json

query.py build validates and rebuilds SQLite.
Runtime tools open the database read-only.

Each claim keeps its period, unit, geography, status, source ID, URL, locator and caveats.

Compatible comparisons use Python arithmetic.
Missing coverage stays explicit.

A checked market comparison

7.4% to 7.2%

–0.2 pp / –20 bps

Central London vacancy
Q1 and Q2 2026
F01/S01 and F02/S02

Q3 is missing in this fixture.
A Q2–Q3 change is unsupported.

Reproduction and verification

A new checkout can run the core checks

After provider setup in the installation guide:

```
python3 scripts/install.py --profile core
. .venv/bin/activate
python scripts/verify_python_cp6.py
```

Optional Apple model profile

```
python scripts/install.py --profile apple
python apple_provider.py
python agent.py --provider apple --interactive
```

Recorded live evaluation

18 / 18

answered prompts
6 conversations
20 tool calls, 0 errors

Chosen Apple profile only.
Automated evidence checks
plus completed human review.